

Data Source Shortlist

Eurovision Voting Blocs & Song Features, 2015 – 2025 · Marta Prandin · Session 1, Deliverable 3

Three candidate sources, ranked by feasibility. For each: official channel, robots / ToS check, expected fields, sample size, and a short verdict. Strategy follows the Slide 12 heuristic — open data file → API → scraping last.

S1 Kaggle — Eurovision Song Contest scores 1975-2024 (Spijkervet)

Open data file (CSV) · kaggle.com/datasets/diamondsnake/eurovision-song-contest-scores-19752024

VERDICT — USE FIRST. Direct CSV download, no auth, license = CC-BY-4.0.

What I get:

- Per-entry results: country, year, song title, performer, language, jury_points, televote_points, total_points, final_rank.
- Voting matrix (~11k rows): voter_country × performer_country × year × jury/televote points.
- Coverage 1975–2024. Ten years of finals (2015–2024) is exactly the window I need; 2025 will be patched in from S3.

Feasibility checks:

- API docs / ToS read: Kaggle CC-BY-4.0 — share-alike with attribution. [OK]
- robots.txt: not relevant (downloadable file, not scraped). [OK]
- Volume: ~30 MB total, well under any limit. [OK]
- GDPR: contains country names, not individuals. [OK]
- Risks: missing tempo/BPM and key (audio features) — handled by S2 in the canvas as Spotify lookup, but the headline analysis works without them.

S2 Wikipedia — “Eurovision Song Contest YYYY” pages

Scrape (pandas.read_html on wikitable) · en.wikipedia.org/wiki/Eurovision_Song_Contest_2025 (one page per year)

VERDICT — USE TO FILL GAPS. CC-BY-SA, structured tables, robots.txt allows it.

What I get:

- Per-year tables: artist, song, language, songwriters, running order, semifinal points and grand-final points.
- 2025 results (the year missing from S1).
- Cross-check column for languages and running order against S1.

Feasibility checks:

- ToS: text is CC-BY-SA-3.0 — attribution + share-alike required. [OK]
- robots.txt: en.wikipedia.org/robots.txt allows User-agent * with crawl-delay; pandas.read_html on a single page per year is well within the spirit of the policy. [OK]
- Rate: 11 pages total (2015–2025), one fetch each — no throttling concerns.
- Volume: ~2 MB total HTML. [OK]
- GDPR: artists are public figures performing under stage names — out of scope of GDPR personal-data concerns. [OK]
- Risks: wikitable layout can change between years; defensive parsing with try/except per year is cheap.

S3 eurovision.tv — official results pages

Scrape (BeautifulSoup) — fallback only · eurovision.tv/event/<year>/grand-final

VERDICT — USE LAST. Site is heavy and JS-rendered for some panels; ToS forbids commercial reuse.

What I get:

- Authoritative full per-country jury and televote score breakdowns (12, 10, 8...1 points given).
- Performance order, presenter info, host city, venue capacity (nice-to-have for a context slide).

Feasibility checks:

- ToS: eurovision.tv/terms — content is for personal, non-commercial use. Academic use with attribution is acceptable for a course project; will not be redistributed. [CAUTION]
- robots.txt: eurovision.tv/robots.txt allows /event/ paths; disallows /admin/ and /api/. [OK]
- Rate: keep ≤ 1 request / 2s, identify with a real User-Agent string.
- Volume: $\sim 11 \text{ pages} \times \sim 300 \text{ KB} = 3 \text{ MB}$. [OK]
- Risks: JS-rendered tables in some years may need playwright instead of requests + bs4 — plan B is to scrape from Wikipedia (S2) and accept loss of some televote-detail granularity.
- Decision: only used if S1 + S2 leave gaps in the per-pair voting data for 2024–2025.

Acquisition strategy

Order of operations for Sessions 2–3:

- Download S1 CSVs into data/raw/ on day one — sacred, never overwritten.
- For 2025, scrape S2 once and save raw HTML alongside the parsed CSV in data/raw/2025_wikipedia.html.
- Only if a 2024–2025 voting-pair row is missing, fall back to S3 with rate-limiting and User-Agent identification.
- Citations file: data/raw/SOURCES.md listing dataset, version, retrieval date, and license per source.

Ethics & reproducibility

- No personal data: all subjects are public-figure performers. No emails, IPs, or addresses are in scope.
- Attribution is mandatory: README and final deck cite Spijkervet (Kaggle), Wikipedia contributors (CC-BY-SA), and EBU/eurovision.tv.
- Reproducibility: every download is dated and pinned by URL in data/raw/SOURCES.md; pip freeze locks the environment.